



Probabilistic Modeling with Tensorflow Made Easy

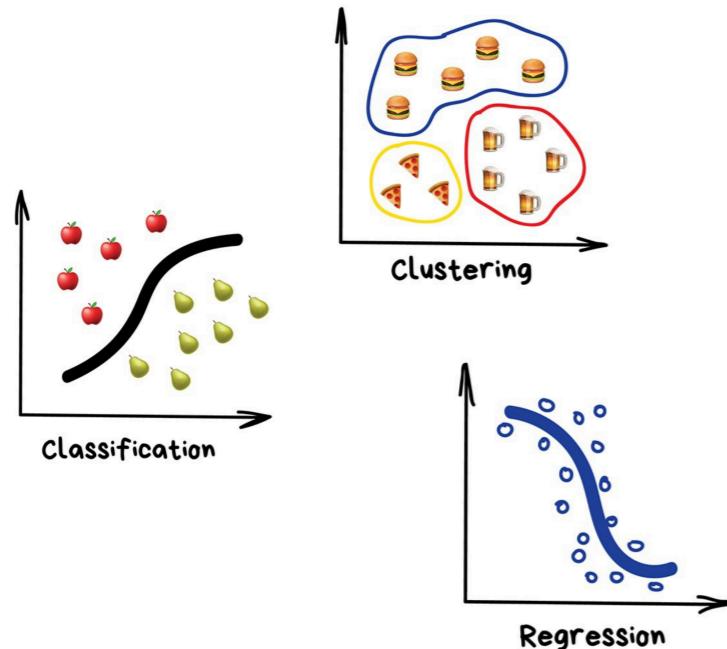
Rafael Cabañas
PIDSIA
Manno, 28.05.19



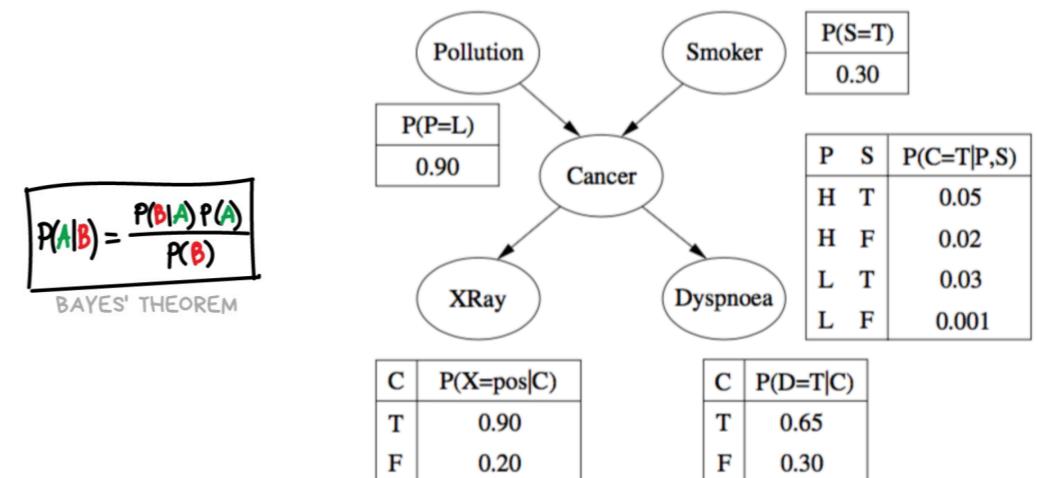
UNIVERSIDAD
DE ALMERÍA



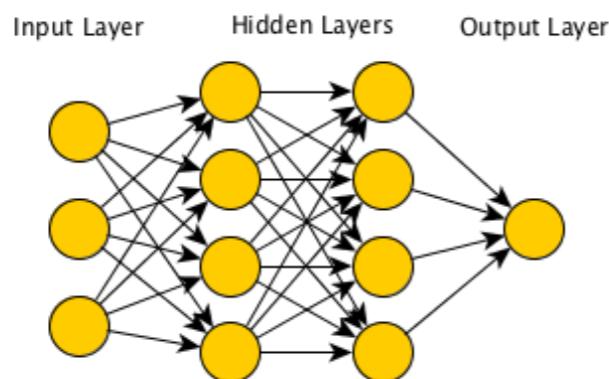
Machine Learning applications



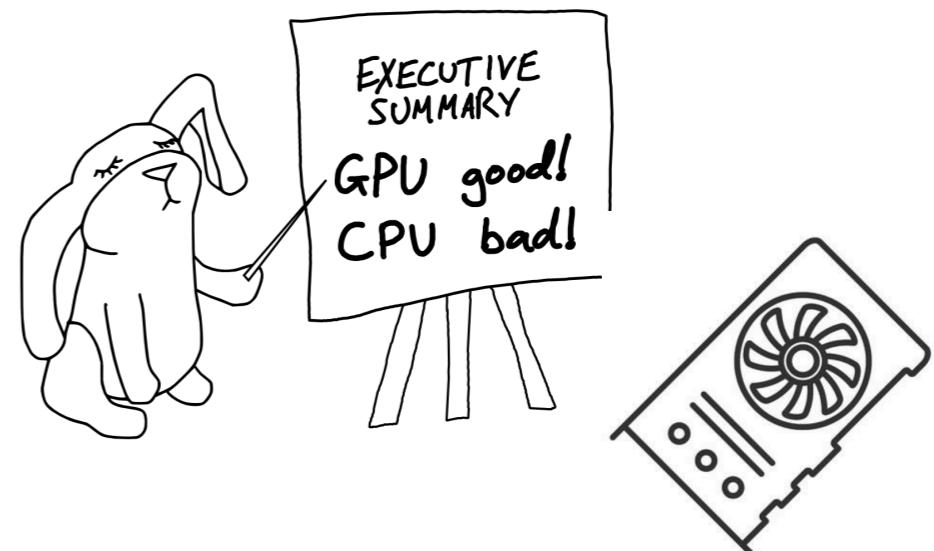
Probabilistic models



Neural networks



Parallelism



Probabilistic programming languages (PPLs)

- PPLs aim to apply the ideas of high-level programming languages to machine learning
- Probabilistic models are defined as programs

```

Sepal.Length Sepal.Width Petal.Length
1          5.1         3.5      1.4
2          4.9         3.0      1.4
3          4.7         3.2      1.3
4          4.6         3.1      1.5
5          5.0         3.6      1.4
6          5.4         3.9      1.7
7          4.6         3.4      1.4
8          5.0         3.4      1.5
9          4.4         2.9      1.4
10         4.9         3.1      1.5
11         5.4         3.7      1.5
12         4.8         3.4      1.6
13         4.8         3.0      1.4
14         4.3         3.0      1.1
15         5.8         4.0      1.2
16         5.7         4.4      1.5
17         5.4         3.9      1.3
18         5.1         3.5      1.4
19         5.7         3.8      1.7
20         5.1         3.8      1.5
21         5.4         3.4      1.7
22         5.1         3.7      1.5
23         4.6         3.6      1.0
24         5.1         3.3      1.7
25         4.8         3.4      1.9

```

```

K, d, N = 5, 10, 200

# model definition
with inf.ProbModel() as m:
    #define the weights
    with inf.replicate(size=K):
        w = inf.models.Normal(0, 1, dim=d)

    # define the generative model
    with inf.replicate(size=N):
        z = inf.models.Normal(0, 1, dim=K)
        x = inf.models.Normal(inf.matmul(z,w),
                             1.0, observed=True, dim=d)

```

data

model as
a program
(samples generator)

 $p(\theta | \text{data})?$ $\mathcal{N}(3.25, 1.2)$

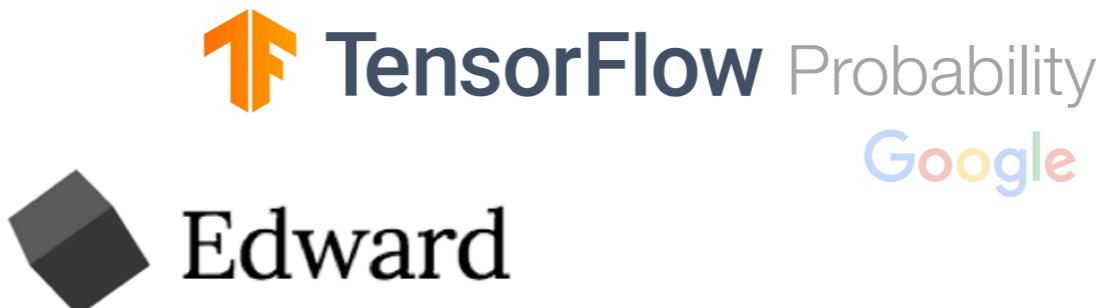
query

blackbox
methods

answer

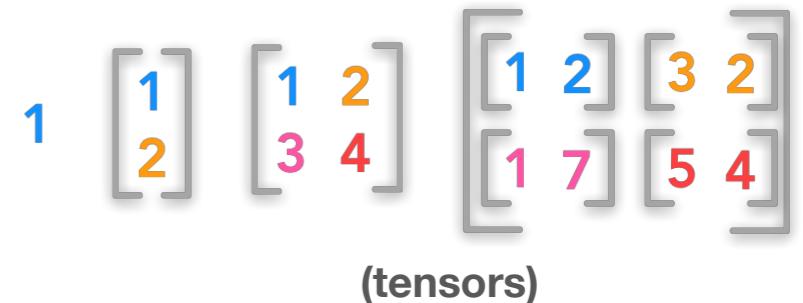
(e.g., Variational Inference)

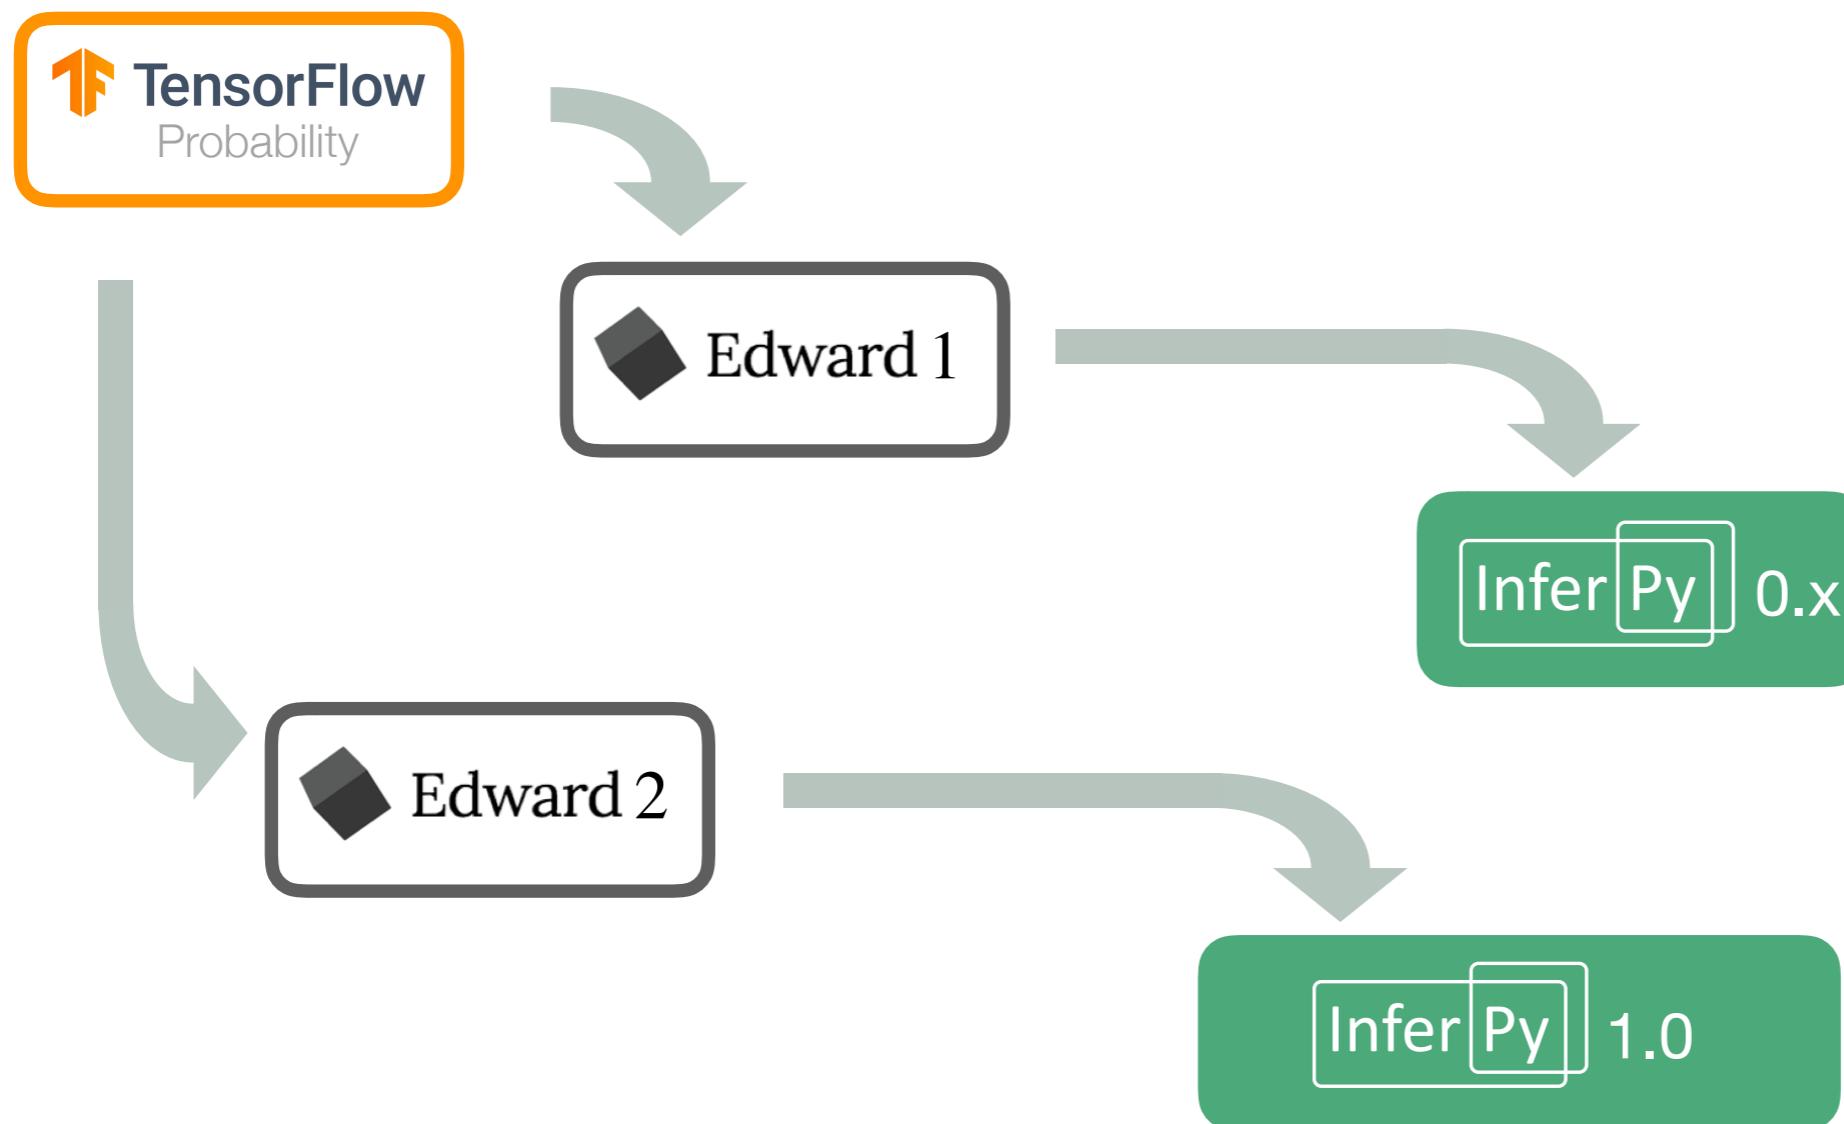
Related software



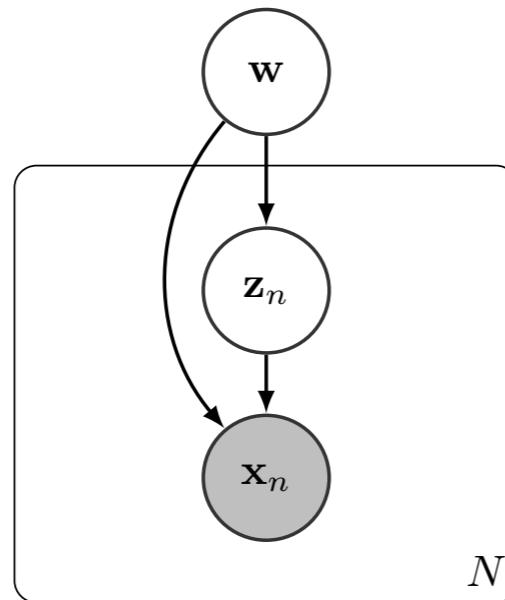
Advantages of InferPy:

- The definition of distributions over tensors is simplified
- No need to have a strong background in inference algorithms





Hierarchical Probabilistic Models



w : global parameters

z : local hidden variables

x : observations

N : number of observations

$p(w)$: prior model

$p(x, z|w)$: data model

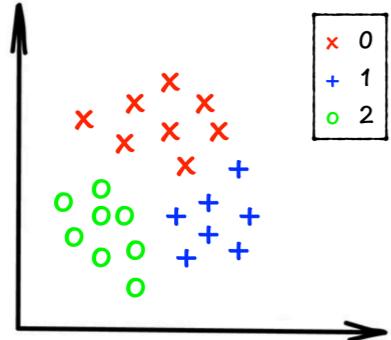
Objective: posterior distribution $p(w, z|x)$

- Dependencies between variables might be defined with TF functions or even NNs



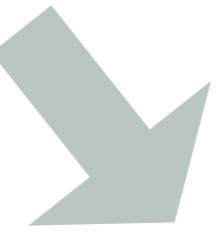
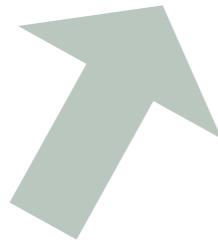
Factor Analysis (PCA)

$k = 2$

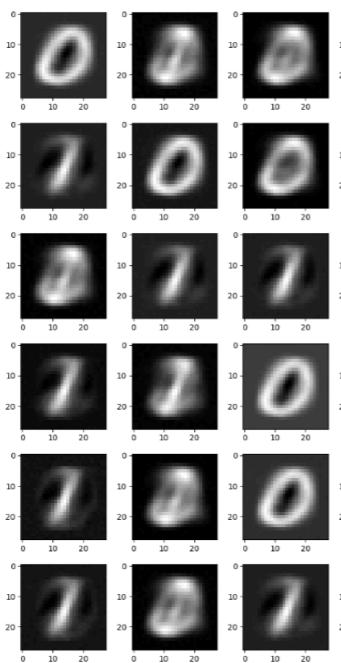
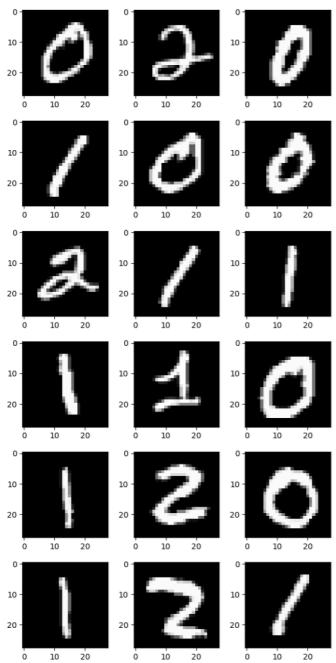


hidden representation

encoding

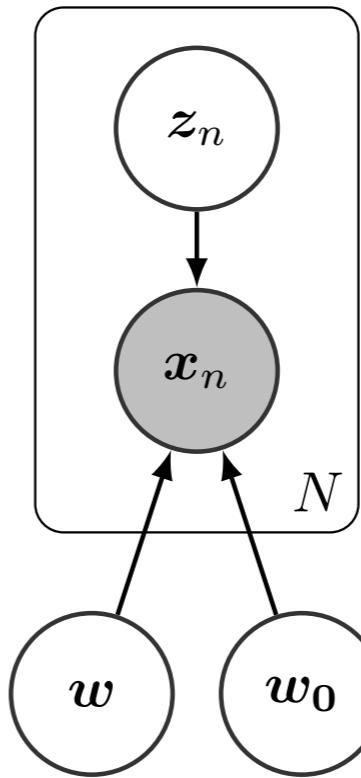


decoding



$$d = 28 * 28 = 784$$

original representation



$$z_n \sim \mathcal{N}_k(0, I)$$

$$x_n \sim \mathcal{N}_d(z_n \cdot w^T + w_0, I)$$

$$w \sim \mathcal{N}_{d \times k}(0, I)$$

$$w_0 \sim \mathcal{N}(0, I)$$

Model definition



```
import inferpy as inf
import tensorflow as tf

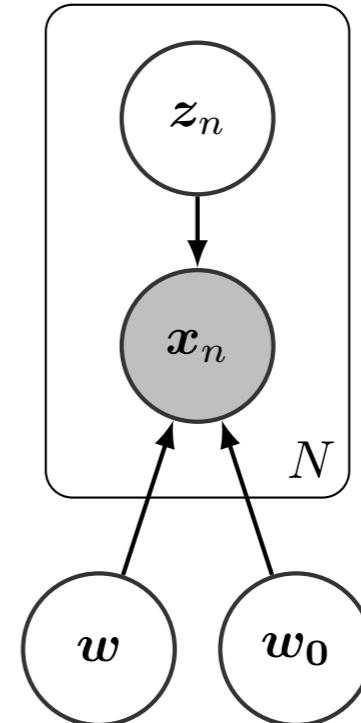
# definition of a generic model
@inf.probmodel
def pca(k,d):
    w = inf.Normal(loc=tf.zeros([k,d]),
                   scale=1, name="w")      # [k,d]

    w0 = inf.Normal(loc=tf.zeros([d]),
                    scale=1, name="w0")     # [d]

    with inf.datamodel():

        z = inf.Normal(tf.zeros([k]),1, name="z") # [N,k]
        x = inf.Normal(z @ w + w0, 1, name="x")   # [N,d]

# create an instance of the model
m = pca(k=2,d=28*28)
```



$$z_n \sim \mathcal{N}_k(0, I)$$

$$x_n \sim \mathcal{N}_d(z_n \cdot w^T + w_0, I)$$

$$w \sim \mathcal{N}_{d \times k}(0, I)$$

$$w_0 \sim \mathcal{N}(0, I)$$

```
In[*]: m.log_prob(m.sample())
```

```
Out[*]:
```

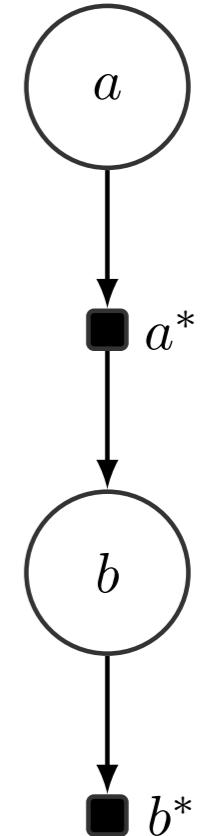
```
{'w': array([-0.96455336, -1.51091952, ..., -1.37483748]),
 'w0': array([-0.84543316]),
 'z': array([-2.04824067, ..., -0.98934147]),
 'x': array([-99.06889862, -16.68617233, ..., -1.40110224])}
```

- Random variables in InferPy encapsulate another from Edward 2
- These can follow the following distributions

```
In[*]: inf.models.random_variable.distributions_all  
Out[*]: ['Autoregressive', 'Bernoulli', 'Beta', 'Binomial',  
'Categorical', 'Cauchy', 'Chi2', 'ConditionalTransformedDistribution',  
'Deterministic', 'Dirichlet', 'DirichletMultinomial',  
'ExpRelaxedOneHotCategorical', 'Exponential', 'Gamma', 'Geometric',  
'HalfNormal', 'Independent', 'InverseGamma', 'Kumaraswamy', 'Laplace',  
'Logistic', 'Mixture', 'MixtureSameFamily', 'Multinomial',  
'MultivariateNormalDiag', 'MultivariateNormalFullCovariance',  
'MultivariateNormalTriL', 'NegativeBinomial', 'Normal',  
'OneHotCategorical', 'Poisson', 'PoissonLogNormalQuadratureCompound',  
'QuantizedDistribution', 'RelaxedBernoulli', 'RelaxedOneHotCategorical',  
'SinhArcsinh', 'StudentT', 'TransformedDistribution', 'Uniform',  
'VectorDeterministic', 'VectorDiffeomixture', 'VectorExponentialDiag',  
'VectorLaplaceDiag', 'VectorSinhArcsinhDiag', 'Wishart']
```

```
a = inf.Normal(0, 100)  
b = inf.Normal(a, 5)
```

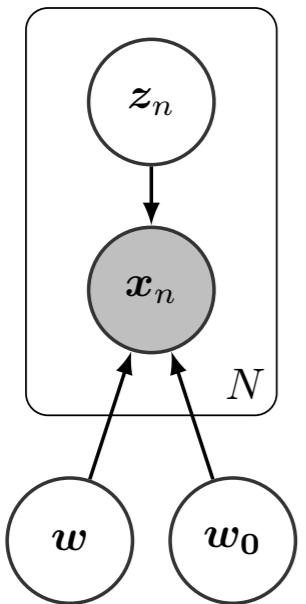
```
In [*]:  
....: sess = inf.get_session()  
....: for i in range(5):  
....:     print(sess.run([a,b]))  
  
[-7.2810316, -6.471646]  
[29.092255, 37.471718]  
[74.87469, 62.43242]  
[44.46464, 39.6697]  
[169.10535, 173.74834]
```



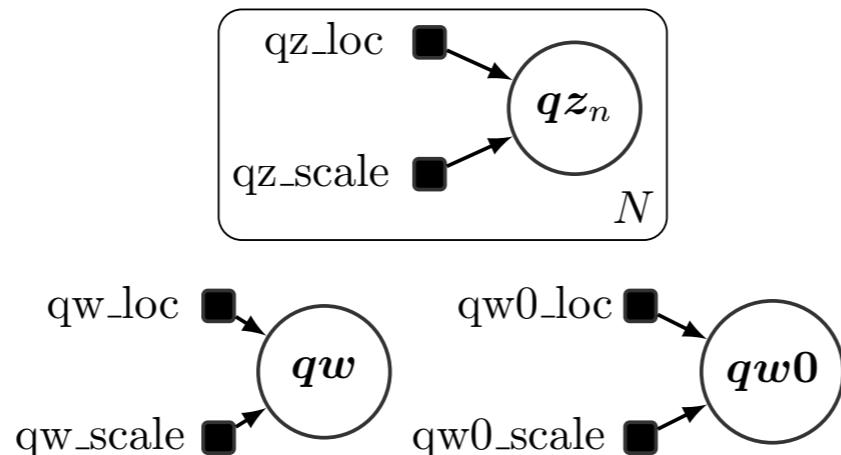
```
# a continuous variable might be parent of a discrete one  
x = inf.Normal(0, 1)  
c = inf.Categorical(probs=tf.case({ x > 0: lambda : [0.0, 1.0],  
                                     x <= 0: lambda : [1.0, 0.0]}))
```

Variational Inference (VI)

P model



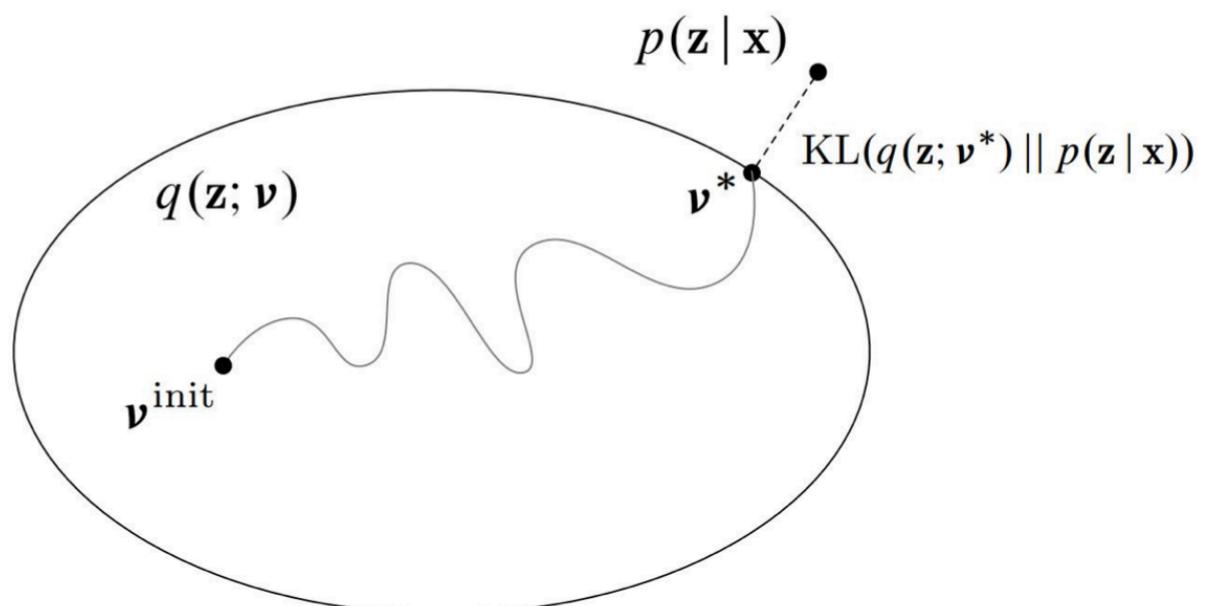
Q model



data

	Sepal.Length	Sepal.Width	Petal.Length
1	5.1	3.5	1.4
2	4.9	3.0	1.4
3	4.7	3.2	1.3
4	4.6	3.1	1.5
5	5.0	3.6	1.4
6	5.4	3.9	1.7
7	4.6	3.4	1.4
8	5.0	3.4	1.5
9	4.4	2.9	1.4
10	4.9	3.1	1.5
11	5.4	3.7	1.5
12	4.8	3.4	1.6
13	4.8	3.0	1.4
14	4.3	3.0	1.1
15	5.8	4.0	1.2
16	5.7	4.4	1.5
17	5.4	3.9	1.3
18	5.1	3.5	1.4
19	5.7	3.8	1.7
20	5.1	3.8	1.5
21	5.4	3.4	1.7
22	5.1	3.7	1.5
23	4.6	3.6	1.0
24	5.1	3.3	1.7
25	4.8	3.4	1.9

- Inference turns into an optimisation problem

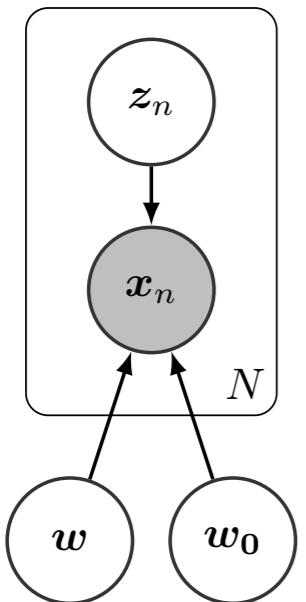


Inference (of the parameters)

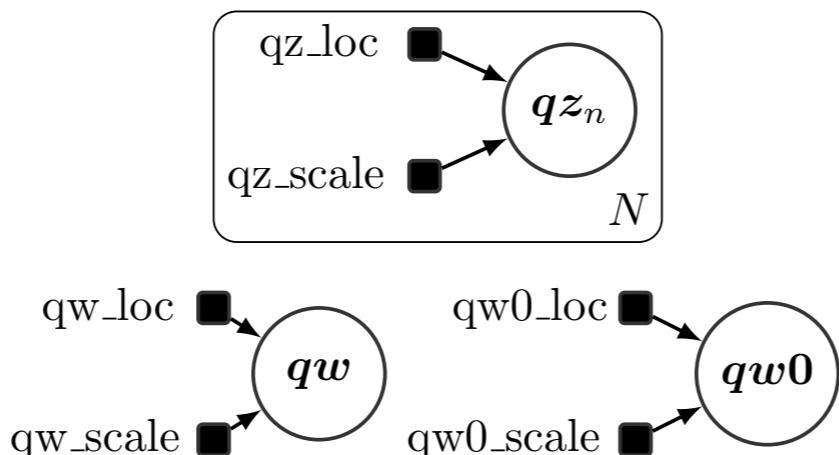


Variational Inference (VI)

P model

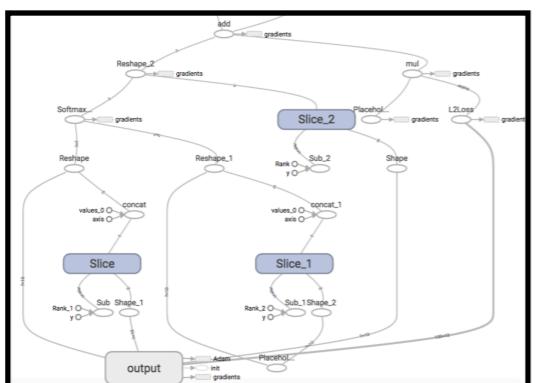


Q model



data

	Sepal.Length	Sepal.Width	Petal.Length
1	5.1	3.5	1.4
2	4.9	3.0	1.4
3	4.7	3.2	1.3
4	4.6	3.1	1.5
5	5.0	3.6	1.4
6	5.4	3.9	1.7
7	4.6	3.4	1.4
8	5.0	3.4	1.5
9	4.4	2.9	1.4
10	4.9	3.1	1.5
11	5.4	3.7	1.5
12	4.8	3.4	1.6
13	4.8	3.0	1.4
14	4.3	3.0	1.1
15	5.8	4.0	1.2
16	5.7	4.4	1.5
17	5.4	3.9	1.3
18	5.1	3.5	1.4
19	5.7	3.8	1.7
20	5.1	3.8	1.5
21	5.4	3.4	1.7
22	5.1	3.7	1.5
23	4.6	3.6	1.0
24	5.1	3.3	1.7
25	4.8	3.4	1.9



Computational graph of the *ELBO*

(to maximise)

$$\text{loss}(P, Q, \text{data}) = -\text{ELBO}(P, Q, \text{data}) \quad (\text{to minimise})$$

Inference (of the parameters)



```
@inf.probmodel
def qmodel(k,d):

    qw_loc = inf.Parameter(tf.zeros([k,d]), name="qw_loc")
    qw_scale = tf.math.softplus(inf.Parameter(tf.ones([k,d]), name="qw_scale"))
    qw = inf.Normal(qw_loc, qw_scale, name="w")

    qw0_loc = inf.Parameter(tf.ones([d]), name="qw0_loc")
    qw0_scale = tf.math.softplus(inf.Parameter(tf.ones([d]), name="qw0_scale"))
    qw0 = inf.Normal(qw0_loc, qw0_scale, name="w0")

    with inf.datamodel():

        qz_loc = inf.Parameter(np.zeros([k]), name="qz_loc")
        qz_scale = tf.math.softplus(inf.Parameter(tf.ones([k]), name="qz_scale"))
        qz = inf.Normal(qz_loc, qz_scale, name="z")

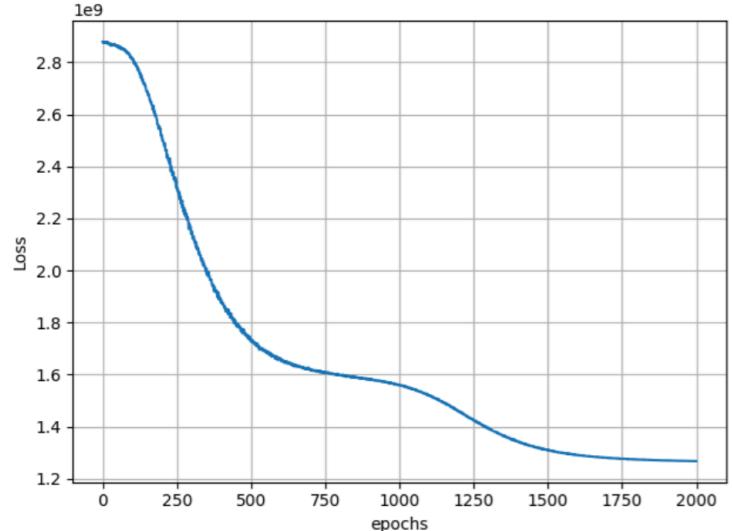
# instance of the Q model
q = qmodel(k=2,d=28*28)

# inference
optimizer = tf.train.AdamOptimizer(learning_rate=0.01)

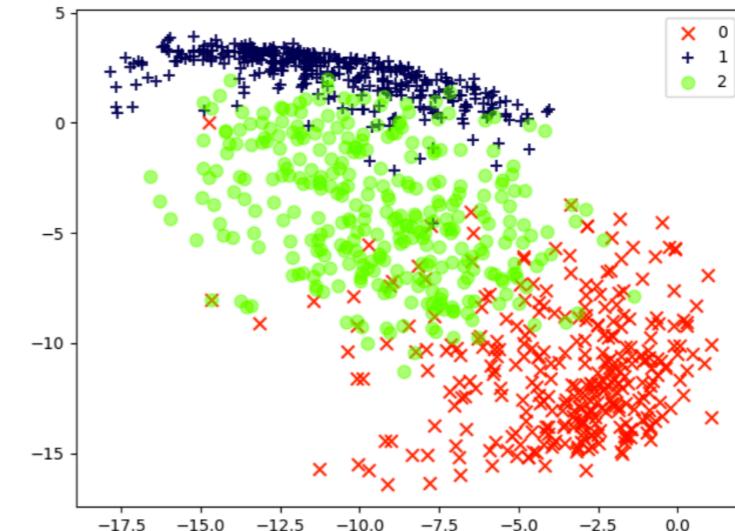
VI = inf.inference.VI(q, optimizer=optimizer, epochs=2000)

m.fit({"x": x_train}, VI)
```

After the inference

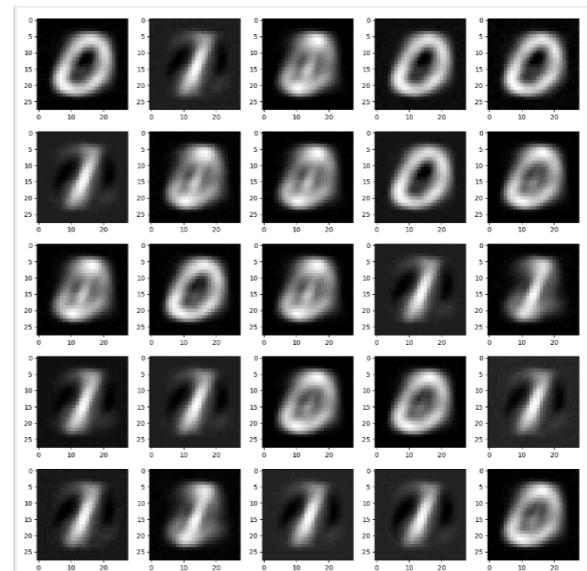


```
# extract the loss evolution  
L = VI.losses
```

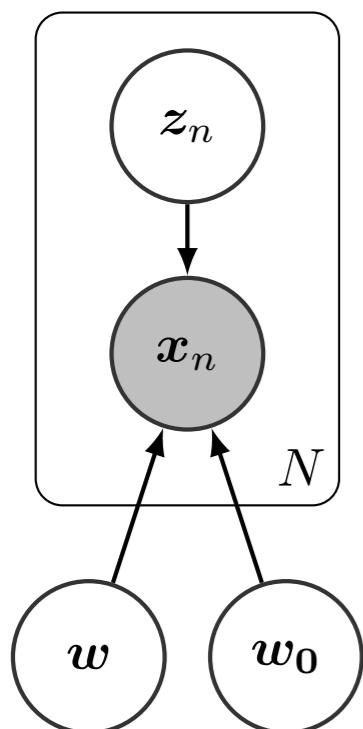


```
# extract the hidden encoding  
sess = inf.get_session()  
sess.run(m.posterior["z"].loc)
```

```
post_samples = m.post_predictive_sample()  
mnist.plot_digits(post_samples["x"], [5,5])
```



Linear PCA

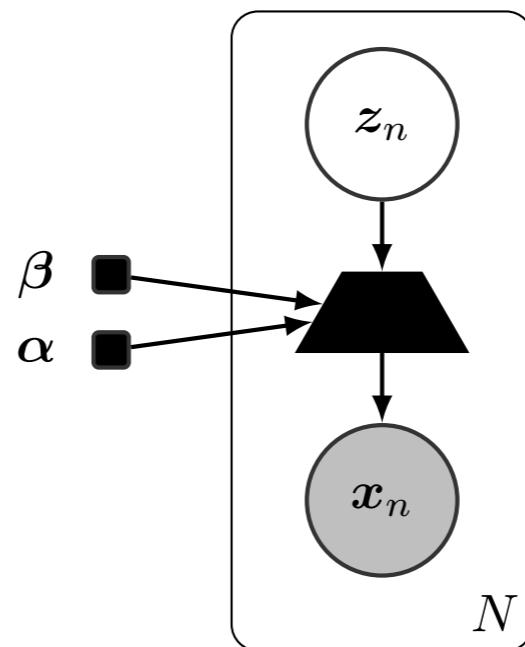


$$z_n \sim \mathcal{N}_k(0, I)$$

$$x_n \sim \mathcal{N}_d(z_n \cdot w^T + w_0, I)$$

$$\begin{aligned} w &\sim \mathcal{N}_{d \times k}(0, I) \\ w_0 &\sim \mathcal{N}(0, I) \end{aligned}$$

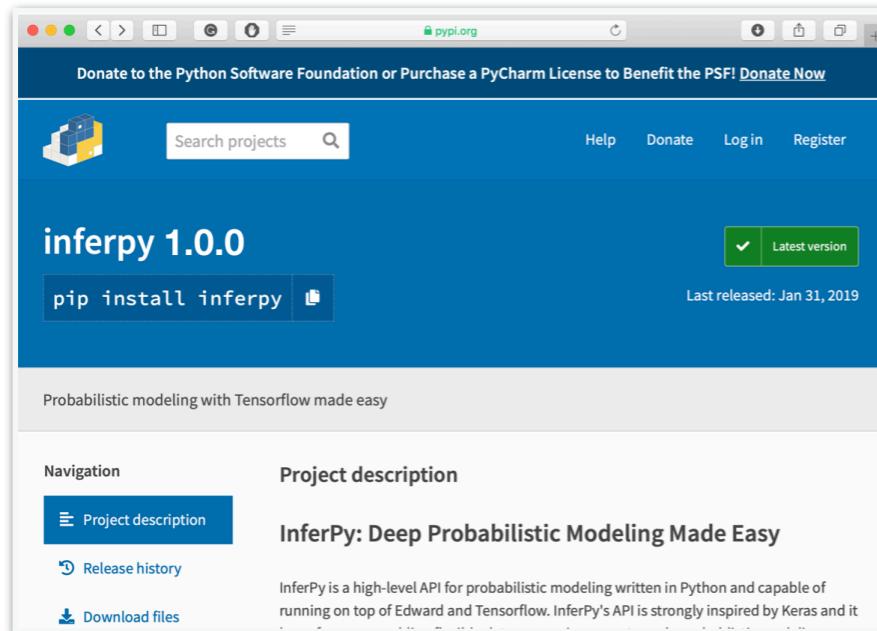
Non-Linear PCA



$$z_n \sim \mathcal{N}_k(0, I)$$

$$x_n \sim \mathcal{N}_d(f(z_n : \beta, \alpha), I)$$

- The same inference methods can be used
- Demo in jupyter notebook



<https://pypi.org/project/inferpy/>

```
$ pip install inferpy
```



The screenshot shows the InferPy Documentation page at version 0.2.2. The left sidebar contains links for QUICK START, Getting Started, Gating Principles, Requirements, GUIDES, MODEL ZOO, PACKAGE REFERENCE, and INFERENCE. The main content area is titled "InferPy: Probabilistic Modeling with Tensorflow Made Easy" and features a large "InferPy" logo. Below the logo, there is a brief introduction and a bulleted list of reasons to use InferPy.

<https://inferpy.readthedocs.io>

The screenshot shows the GitHub repository page for "PGM-Lab / InferPy". It displays basic repository statistics: 354 commits, 8 branches, 9 releases, and 2 contributors. The "Code" tab is selected, showing a list of recent commits. Commits include updates to Dockerfiles, documentation, and bug fixes.

<https://github.com/PGM-Lab/InferPy>

The screenshot shows the full-text article "InferPy: Probabilistic modeling with Tensorflow made easy" by Cabañas, R., Salmerón, A., & Masegosa, A. R. (2019). The article is published in the journal "Knowledge-Based Systems" (Volume 168, pages 25-27). It includes sections for Article History, Abstract, and Keywords.

Cabañas, R., Salmerón, A., & Masegosa, A. R. (2019).
InferPy: Probabilistic modeling with Tensorflow made easy. *Knowledge-Based Systems*, 168, 25-27.